

Automatic Detection of Whale Lunges: A Deep Learning Approach

Will McCloskey
Hugo Valdivia

Roadmap

- Background + Problem Definition
- Dataset
- Results
- Algorithm Details
- Neural Networks Overview

Background

- Many species of whales engage in lunge-feeding behavior
- Analysis is important for understanding the ecology, conservation, and evolution of these species
- Currently lunge feeding times are labeled by hand

We want to improve this process

Objective

Automate the labeling process!



Problem Statement

Data:

$\{x^{(t)} : t \in \{1, 2, \dots, T\}\}$ Deployment time series from accelerometer measurements

$y^{(t)} \in \{0, 1\}$ Hand-labeled lunge times

Goal:

Given an unlabeled deployment, predict the lunge times.

Dataset

Species	Blue Whale	Minke Whale
Number of Deployments	29	6
Number of Lunges	4768	6151
Hours of Data	319.7	127.3
Sampling Rate	10 Hz	10 Hz

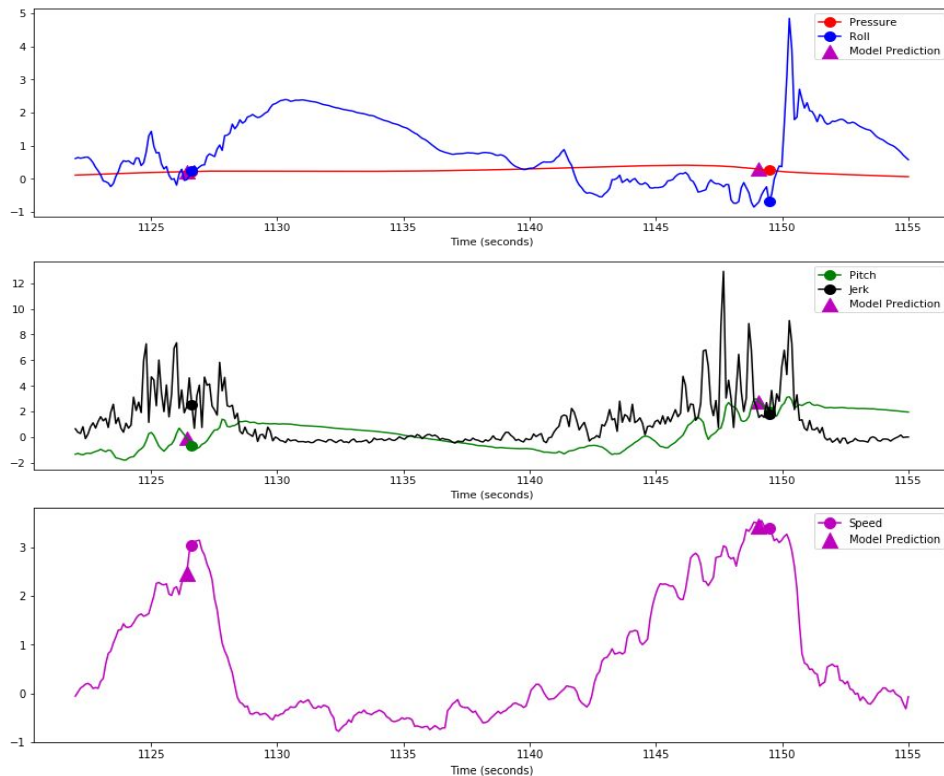
Feature Selection And Preprocessing

- **Features Used**
 - Speed
 - Pitch
 - Jerk
 - Roll
 - Pressure (proxy for depth)
- **Preprocessing**
 - Time series are normalized (mean=0, standard deviation=1) per deployment

Generalization: Key Goal of Machine Learning

- Want model to perform well on unseen data
- To achieve this, split deployments into 3 sets: training, validation, test
 - Model will see the training set but not the others.
- Training set (~80% of lunges)
 - Learning algorithm needs data to learn how to predict
- Validation set (~10% of lunges)
 - Fake test set to prevent overfitting to training set and tune accordingly
- Test set (~10% of lunges)
 - Real world performance test to assess model at the very end

Example Output (Minke Val Deployment)



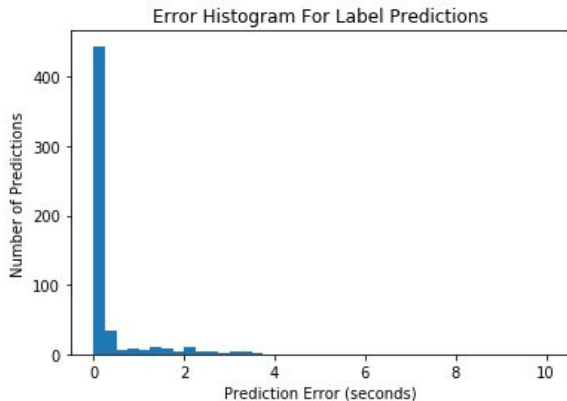
Metrics

- A predicted lunge is considered correct if it is within 5 seconds of a true lunge
 - True lunge = expert-labeled lunge
- True positive rate:
 - Rate at which true lunges are detected
$$\frac{\text{\#true lunges detected}}{\text{\#true lunges}}$$
- False positive rate:
 - Rate at which false predictions are made
$$\frac{\text{\#incorrect predictions}}{\text{\#total number of predictions}}$$
- Number of correct predictions
- Number of predictions
- Number of overcounted correct predictions
- Average prediction error (in seconds)
- F₁ and F₂ scores

Test Set Results (Blue Whales)


Feed Forward Model

The true positive rate for tolerance 5 seconds is 0.97
The false positive rate for tolerance 5 seconds is 0.037
Total lunges in files: 568
Num correct lunges: 551
Num predicted lunges: 572
We overcount by 0
We are off by an average of this many seconds: 0.31724137931034496
The f₁ score is 0.967
The f₂ score is 0.969



ResNet Model

The true positive rate for tolerance 5 seconds is 0.972
The false positive rate for tolerance 5 seconds is 0.03
Total lunges in files: 568
Num correct lunges: 552
Num predicted lunges: 569
We overcount by 0
We are off by an average of this many seconds: 0.48550724637681136
The f₁ score is 0.971
The f₂ score is 0.971




Test Set Results (Minke Whales)

Feed Forward Model

The true positive rate for tolerance 5 seconds is 0.962
The false positive rate for tolerance 5 seconds is 0.088
Total lunges in files: 799
Num correct lunges: 769
Num predicted lunges: 843
We overcount by 0
We are off by an average of this many seconds: 0.14395318595578616
The f_1 score is 0.937
The f_2 score is 0.952


Error Histogram For Label Predictions



ResNet Model

The true positive rate for tolerance 5 seconds is 0.935
The false positive rate for tolerance 5 seconds is 0.141
Total lunges in files: 799
Num correct lunges: 747
Num predicted lunges: 870
We overcount by 3
We are off by an average of this many seconds: 0.21311914323962408
The f_1 score is 0.895
The f_2 score is 0.919

Error Histogram For Label Predictions



Clicks Saved (Test Set)

$\text{\#clicks still needed} = \text{\#incorrect predictions} + \text{\#missed true lunges}$

$\text{\#clicks saved} = \text{\#true lunges} - \text{\#clicks still needed}$

$\text{clicks saved rate} = \frac{\text{\#clicks saved}}{\text{\#true lunges}}$

Blue Whales (568 True Lunges)

Model	Feed Forward	ResNet
Clicks Saved	531	535
Clicks Saved Rate	93.4%	94.1%

Clicks Saved (Test Set)

$\text{\#clicks still needed} = \text{\#incorrect predictions} + \text{\#missed true lunges}$

$\text{\#clicks saved} = \text{\#true lunges} - \text{\#clicks still needed}$

$\text{clicks saved rate} = \frac{\text{\#clicks saved}}{\text{\#true lunges}}$

Minke Whales (799 True Lunges)

Model	Feed Forward	ResNet
Clicks Saved	695	624
Clicks Saved Rate	87.5%	78.1%


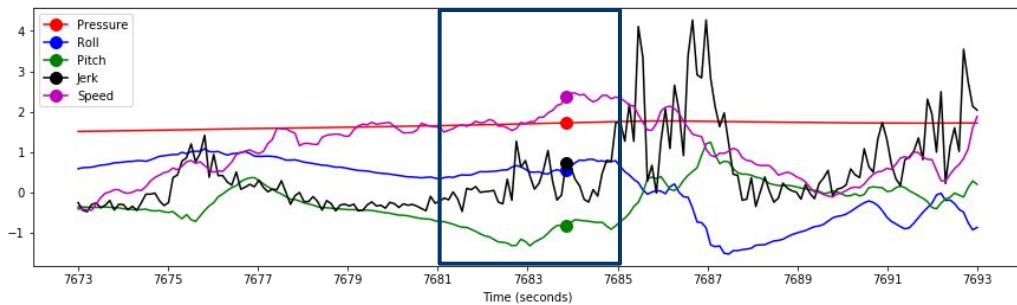
Algorithm Details

Prediction On Windows

- Take a window of length T from the data
 - Assign the window a 1 if there is a lunge within x seconds of the center
 - Assign the window a 0 otherwise

Algorithm Details

Minke Whale Example (20 second windows, 2 second tolerance)



Algorithm Details

Model Prediction

- Model attempts to predict the assigned value (1 or 0)
- Model predicts probability that output is 1

Sliding Windows

- To predict at time t , feed the model the window centered at time t .
- To predict on whole deployment, make predictions every second.

Algorithm Details

Consolidating Predictions

- Output predictions can be noisy
 - Use a moving average filter to reduce false positives
- Now many predictions with probability > 0.5 correspond to a given true lunge
- To consolidate:
 - Cluster output predictions if they are within C seconds of each other
 - Output the max probability prediction (probability P) if $P > T$ where T is a threshold value
 - Values C and T are chosen using performance on deployments in the validation set

Algorithm Details

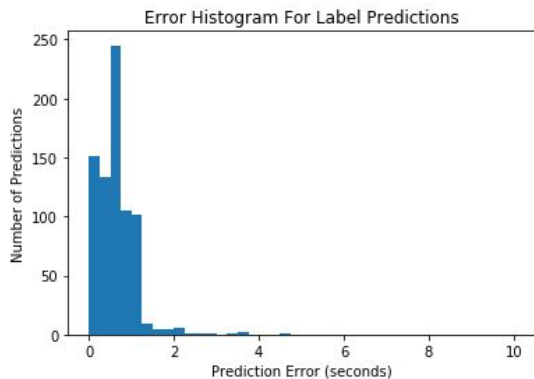
Correcting Predictions

- **First (detection) model detects whether there is a lunge in the middle x seconds**
 - No guarantee that predictions tightly align with true labels
- **Second (correction) model**
 - Takes a predicted time and predicts offset to true label
 - Trained using synthetically-generated incorrect labels
 - Offsets sampled from the distribution of errors that first model created

Correction Model Results (Minke)

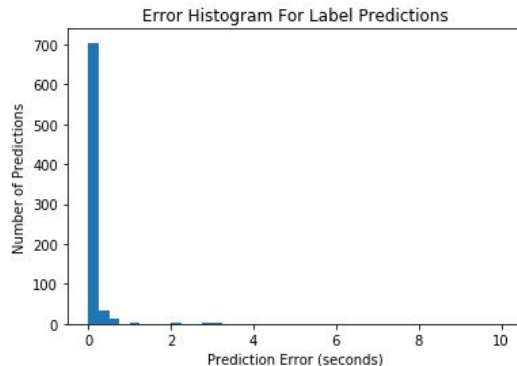
Uncorrected Predictions

The true positive rate for tolerance 5 seconds is 0.962
The false positive rate for tolerance 5 seconds is 0.088
Total lunges in files: 799
Num correct lunges: 769
Num predicted lunges: 843
We overcount by 0
We are off by an average of this many seconds: 0.6243172951885562
The f_1 score is 0.937
The f_2 score is 0.952

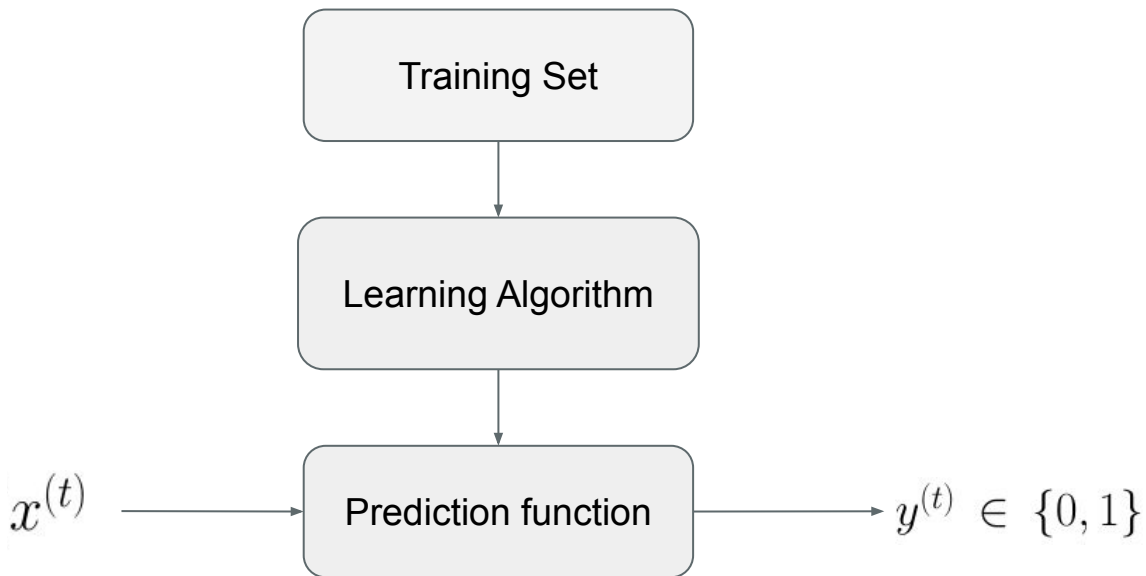


Corrected Predictions



The true positive rate for tolerance 5 seconds is 0.962
The false positive rate for tolerance 5 seconds is 0.088
Total lunges in files: 799
Num correct lunges: 769
Num predicted lunges: 843
We overcount by 0
We are off by an average of this many seconds: 0.14395318595578616
The f_1 score is 0.937
The f_2 score is 0.952



Supervised Learning: Binary Classification



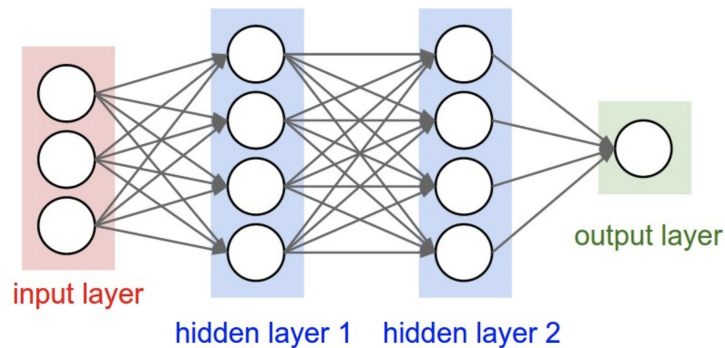
Neural Networks



A cartoon drawing of a biological neuron (left) and its mathematical model (right).


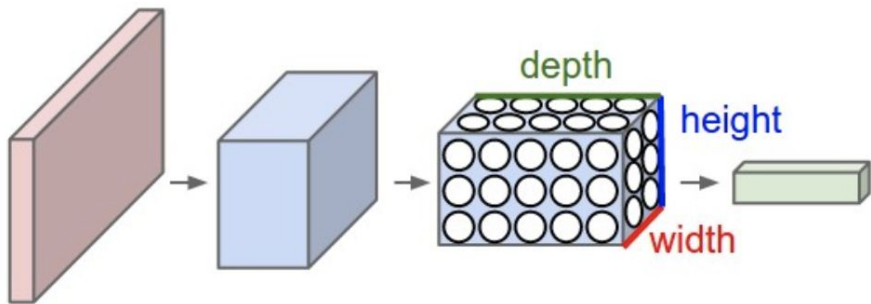
Neural Networks

- Connected, acyclic graph of neurons
- Single neuron linear classifiers
 - Can be equivalent to logistic regression and SVM
- General feed-forward neural networks
 - A series of matrix multiplications linked with non-linearities
- Representational Power
 - Single hidden layer is enough to approximate any continuous function
 - However, just because a function can be learned doesn't mean it will be learned
 - Empirically, deeper networks will work much better in practice but must be tuned appropriately
 - Key takeaway: neural networks are able to learn well-behaving functions that fit real-world data tremendously well.




Convolutional Networks

- Optimized architecture for images
 - Parameter sharing to decrease number of parameters
- Typical architecture consists of CONV, POOL, and FC layers
- Most important is CONV
 - Small filters slide across image to output a new image that holds the filter's response to each spatial section
 - Filters will activate when they detect a certain feature



Convolutional Networks: Visualized




Low-level
features

Mid-level
features

High-level
features


Linearly
separable
classifier



VGG-16 Conv1_1




VGG-16 Conv3_2



VGG-16 Conv5_3

Residual Networks (ResNets)

- Trend towards deeper networks
 - Makes sense to learn more low/mid/high-level features
 - But much harder to train
- ResNets allow for much deeper networks (~1000-layer networks)
- Have been found to work quite well on time series data



External Resources

- Code base can be found at:
 - <https://github.com/valdivia4/Deep-Learning-Lunge-Detection>
- Implemented in Python
- Detailed documentation provided
 - No prior Python experience necessary