

# Automatic Detection of Whale Lunges: A Deep Learning Approach

Will McCloskey  
Hugo Valdivia

# Roadmap

- Background
- Formal Problem Definition
- Dataset
- Results
- Algorithm Details
- Neural Networks Overview

# Background

- Variety of whale species engage in lunge-feeding behavior
- Analysis of such behavior is important for understanding these species
- Hand-labeling of lunge feeding times is inefficient use of time

We want to improve this process.

# Objective

## Automate the labeling process!



# Formal Problem Definition

## Data:

$\{x^{(t)} : t \in \{1, 2, \dots, T\}\}$  Deployment time series from accelerometer measurements

$y^{(t)} \in \{0, 1\}$  Hand-labeled lunge times

## Problem:

Given an unlabeled deployment, predict the lunge times.

# Dataset

Species	Blue Whale	Minke Whale
Number of Deployments	29	6
Number of Lunges	4768	6151
Hours of Data	319.7	127.3
Sampling Rate	10 Hz	10 Hz

# Feature Selection And Preprocessing

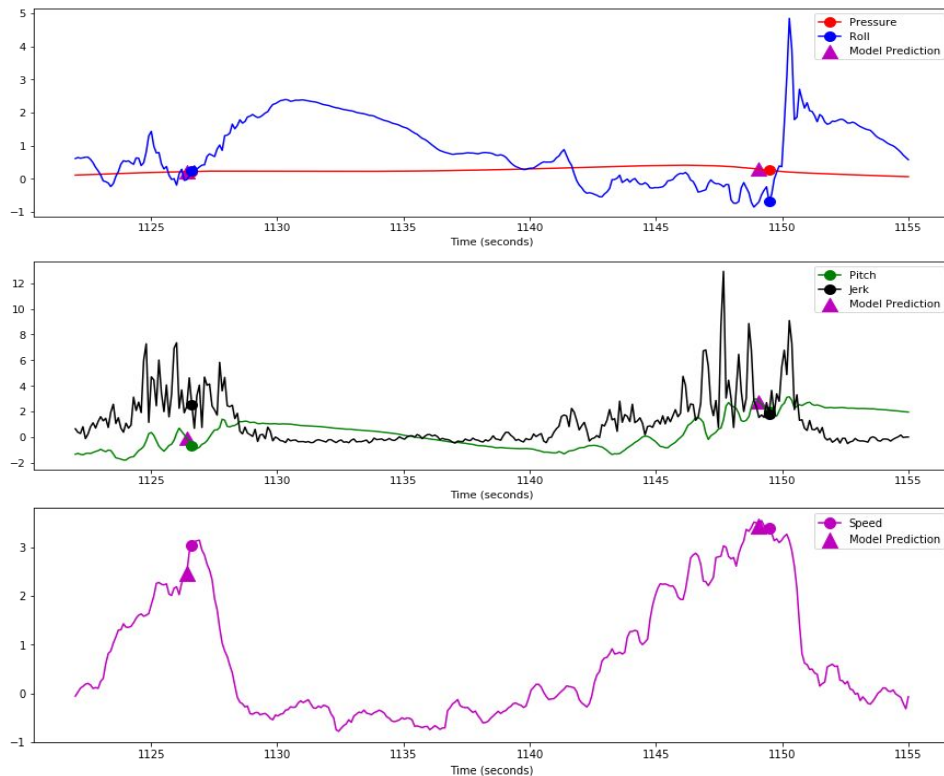
- Features Used
  - Speed
  - Pitch
  - Jerk
  - Roll
  - Pressure (proxy for depth)
- Preprocessing
  - Time series are standardized (mean=0, standard deviation=1) per deployment

# Generalization: Key Goal of Machine Learning

- Want model to perform well on unseen data
- To achieve this, split deployments into 3 sets: training, validation, test
  - Model will see the training set but not the others.
- Training set (~80% of lunges)
  - Learning algorithm needs data to learn how to predict
- Validation set (~10% of lunges)
  - Fake test set to prevent overfitting to training set and tune accordingly
- Test set (~10% of lunges)
  - Real world performance test to assess model at the very end



# Example Output (Minke Val Deployment)



# Metrics

- A predicted lunge is considered correct if it is within 5 seconds of a true lunge
  - True lunge = expert-labeled lunge
- True positive rate:
  - Rate at which true lunges are detected
- False positive rate:
  - Rate at which false predictions are made
- Number of correct predictions
- Number of predictions
- Number of overcounted correct predictions
- Average prediction error (in seconds)
- F<sub>1</sub> and F<sub>2</sub> scores

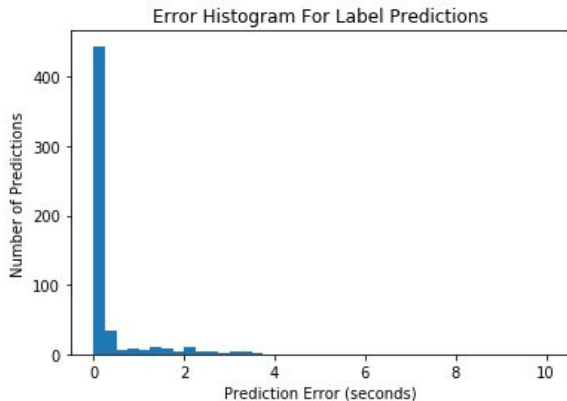
$$\frac{\text{\#true lunges detected}}{\text{\#true lunges}}$$

$$\frac{\text{\#incorrect predictions}}{\text{\#total number of predictions}}$$

# Test Set Results (Blue Whales)

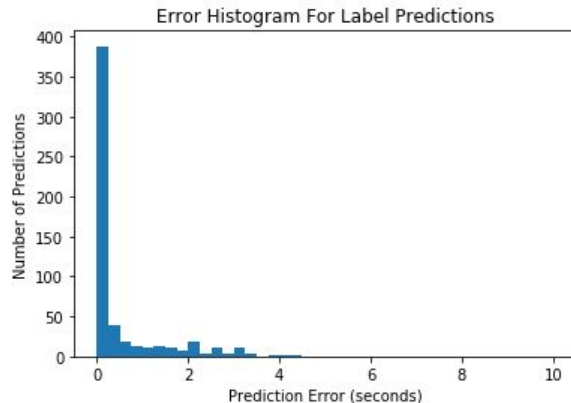
## Feed Forward Model

The true positive rate for tolerance 5 seconds is 0.97  
The false positive rate for tolerance 5 seconds is 0.037  
Total lunges in files: 568  
Num correct lunges: 551  
Num predicted lunges: 572  
We overcount by 0  
We are off by an average of this many seconds: 0.31724137931034496  
The f<sub>1</sub> score is 0.967  
The f<sub>2</sub> score is 0.969



## ResNet Model

The true positive rate for tolerance 5 seconds is 0.972  
The false positive rate for tolerance 5 seconds is 0.03  
Total lunges in files: 568  
Num correct lunges: 552  
Num predicted lunges: 569  
We overcount by 0  
We are off by an average of this many seconds: 0.48550724637681136  
The f<sub>1</sub> score is 0.971  
The f<sub>2</sub> score is 0.971

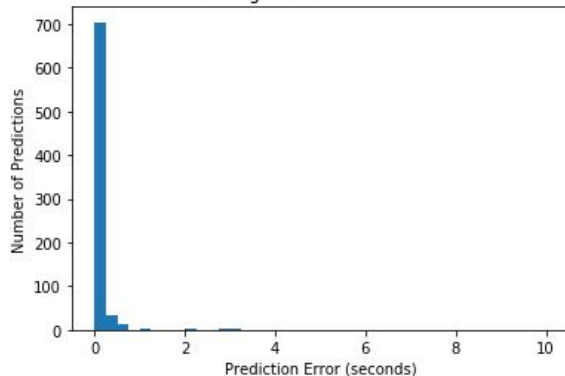


# Test Set Results (Minke Whales)

## Feed Forward Model

The true positive rate for tolerance 5 seconds is 0.962  
The false positive rate for tolerance 5 seconds is 0.088  
Total lunges in files: 799  
Num correct lunges: 769  
Num predicted lunges: 843  
We overcount by 0  
We are off by an average of this many seconds: 0.14395318595578616  
The f\_1 score is 0.937  
The f\_2 score is 0.952

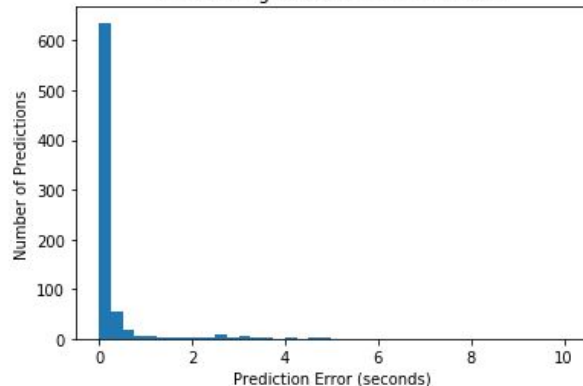
Error Histogram For Label Predictions



## ResNet Model

The true positive rate for tolerance 5 seconds is 0.935  
The false positive rate for tolerance 5 seconds is 0.141  
Total lunges in files: 799  
Num correct lunges: 747  
Num predicted lunges: 870  
We overcount by 3  
We are off by an average of this many seconds: 0.21311914323962408  
The f\_1 score is 0.895  
The f\_2 score is 0.919

Error Histogram For Label Predictions



# Clicks Saved (Test Set)

$\text{\#clicks still needed} = \text{\#incorrect predictions} + \text{\#missed true lunges}$

$\text{\#clicks saved} = \text{\#true lunges} - \text{\#clicks still needed}$

$\text{clicks saved rate} = \frac{\text{\#clicks saved}}{\text{\#true lunges}}$

## Blue Whales (568 True Lunges)

Model	Feed Forward	ResNet
Clicks Saved	531	535
Clicks Saved Rate	93.4%	94.1%

# Clicks Saved (Test Set)

$\text{\#clicks still needed} = \text{\#incorrect predictions} + \text{\#missed true lunges}$

$\text{\#clicks saved} = \text{\#true lunges} - \text{\#clicks still needed}$

$\text{clicks saved rate} = \frac{\text{\#clicks saved}}{\text{\#true lunges}}$

## Minke Whales (799 True Lunges)

Model	Feed Forward	ResNet
Clicks Saved	695	624
Clicks Saved Rate	87.5%	78.1%

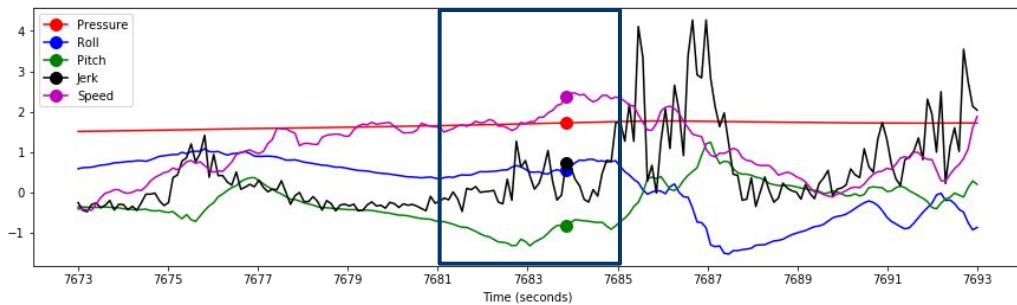
# Algorithm Details

## Prediction On Windows

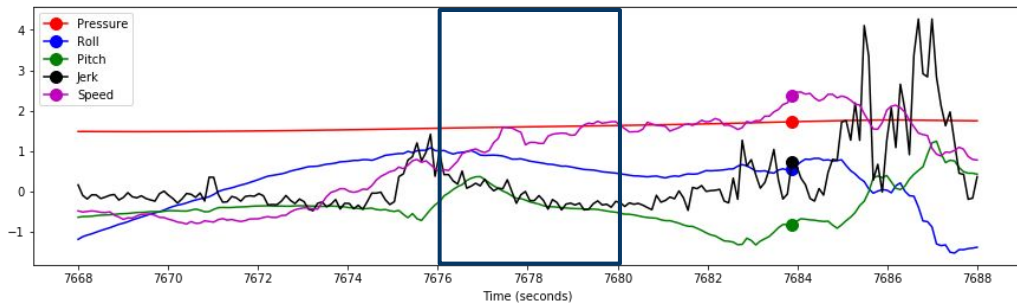
- Take a window of length  $T$  from the data
  - Assign the window a 1 if there is a lunge within  $x$  seconds of the center
  - Assign the window a 0 otherwise

# Algorithm Details

Minke Whale Example (20 second windows, 2 second tolerance)



1



0



# Algorithm Details

## Model Prediction

- Model attempts to predict the assigned value (1 or 0)
- Model predicts probability that output is 1

## Sliding Windows

- To predict at time  $t$ , feed the model the window centered at time  $t$ .
- To predict on whole deployment, make predictions every second.

# Algorithm Details

## Consolidating Predictions

- Output predictions can be noisy
  - Use a moving average filter to reduce false positives
- Now many predictions with probability  $> 0.5$  correspond to a given true lunge
- To consolidate:
  - Cluster output predictions if they are within  $C$  seconds of each other
  - Output the max probability prediction (probability  $P$ ) if  $P > T$  where  $T$  is a threshold value
  - Values  $C$  and  $T$  are chosen using performance on deployments in the validation set

# Algorithm Details

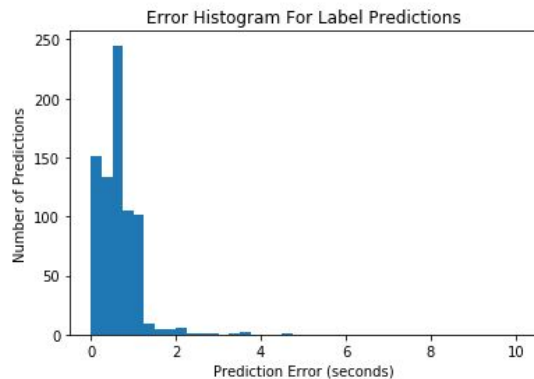
## Correcting Predictions

- First (detection) model detects whether there is a lunge in the middle  $x$  seconds
  - No guarantee that predictions tightly align with true labels
- Second (correction) model
  - Takes a predicted time and predicts offset to true label
  - Trained using synthetically-generated incorrect labels
    - Offsets sampled from the distribution of errors that first model created

# Correction Model Results (Minke)

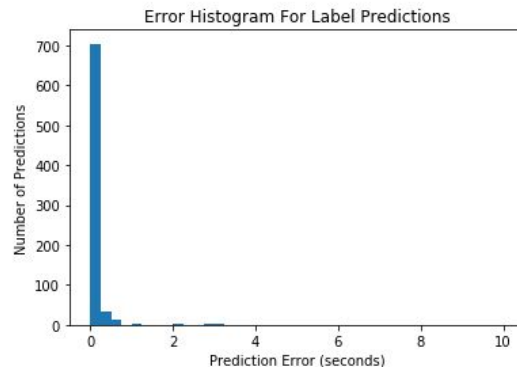
## Uncorrected Predictions

The true positive rate for tolerance 5 seconds is 0.962  
The false positive rate for tolerance 5 seconds is 0.088  
Total lunges in files: 799  
Num correct lunges: 769  
Num predicted lunges: 843  
We overcount by 0  
We are off by an average of this many seconds: 0.6243172951885562  
The f\_1 score is 0.937  
The f\_2 score is 0.952

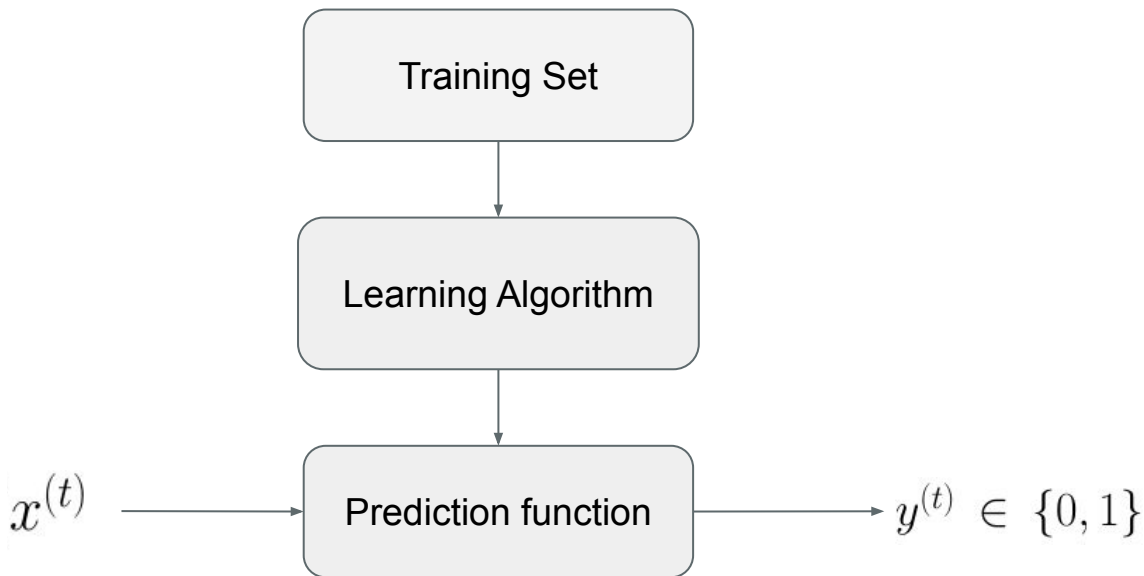


## Corrected Predictions

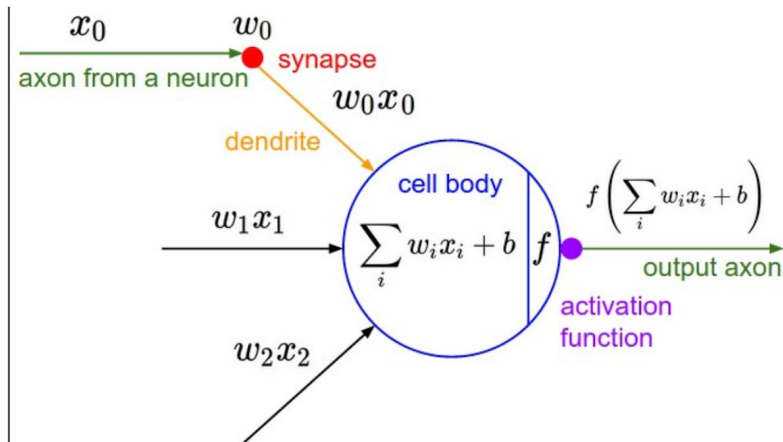
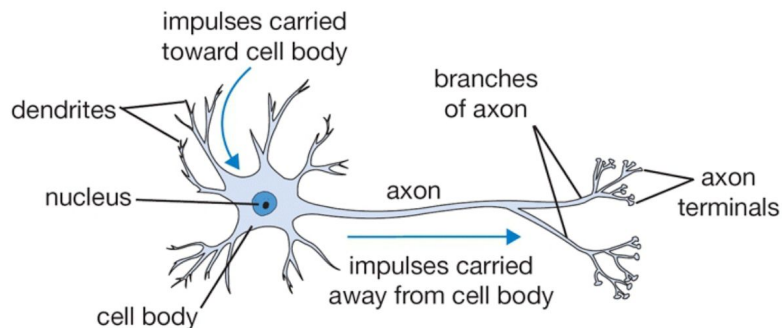
The true positive rate for tolerance 5 seconds is 0.962  
The false positive rate for tolerance 5 seconds is 0.088  
Total lunges in files: 799  
Num correct lunges: 769  
Num predicted lunges: 843  
We overcount by 0  
We are off by an average of this many seconds: 0.14395318595578616  
The f\_1 score is 0.937  
The f\_2 score is 0.952



# Supervised Learning: Binary Classification



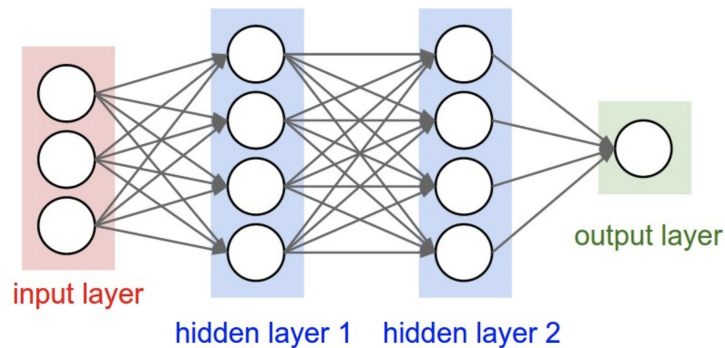
# Neural Networks



A cartoon drawing of a biological neuron (left) and its mathematical model (right).

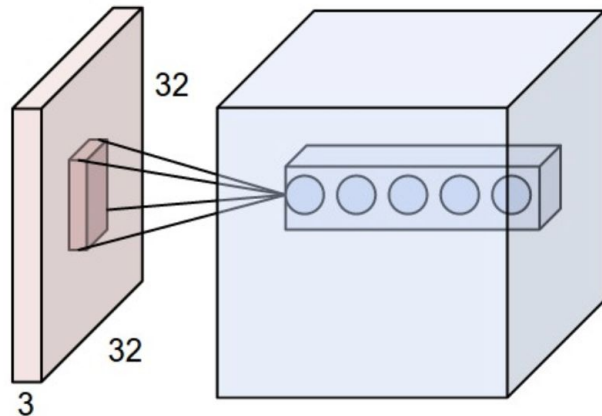
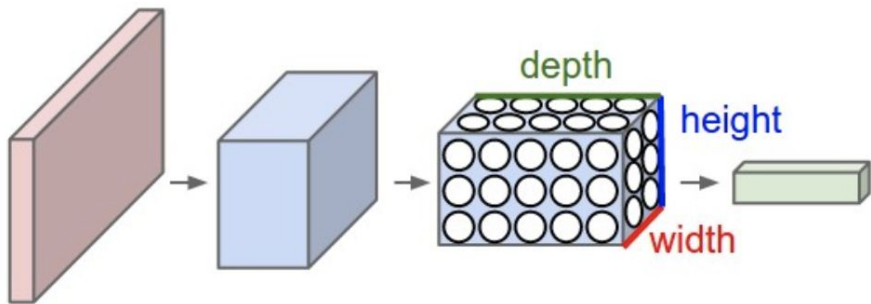
# Neural Networks

- Connected, acyclic graph of neurons
- Single neuron linear classifiers
  - Can be equivalent to logistic regression and SVM
- General feed-forward neural networks
  - A series of matrix multiplications linked with non-linearities
- Representational Power
  - Single hidden layer is enough to approximate any continuous function
  - However, just because a function can be learned doesn't mean it will be learned
    - Empirically, deeper networks will work much better in practice but must be tuned appropriately
  - Key takeaway: neural networks are able to learn well-behaving functions that fit real-world data tremendously well.



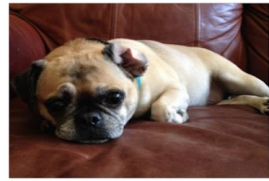
# Convolutional Networks

- Optimized architecture for images
  - Parameter sharing to decrease number of parameters
- Typical architecture consists of CONV, POOL, and FC layers
- Most important is CONV
  - Small filters slide across image to output a new image that holds the filter's response to each spatial section
  - Filters will activate when they detect a certain feature





# Convolutional Networks: Visualized

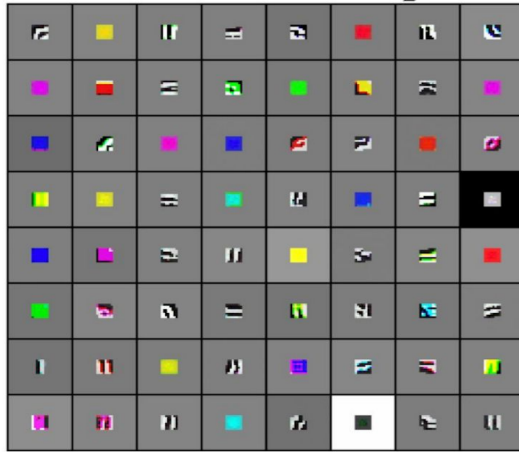


Low-level  
features

Mid-level  
features

High-level  
features

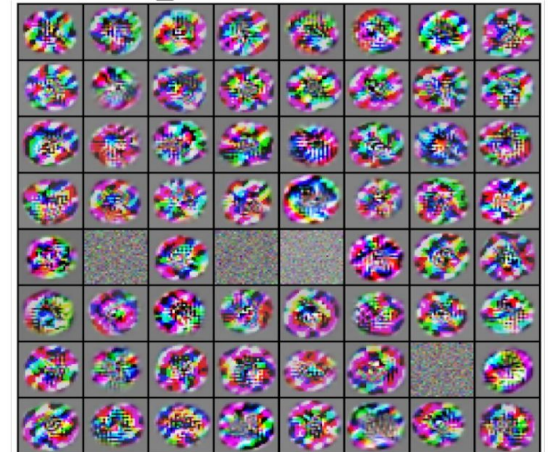
Linearly  
separable  
classifier



VGG-16 Conv1\_1



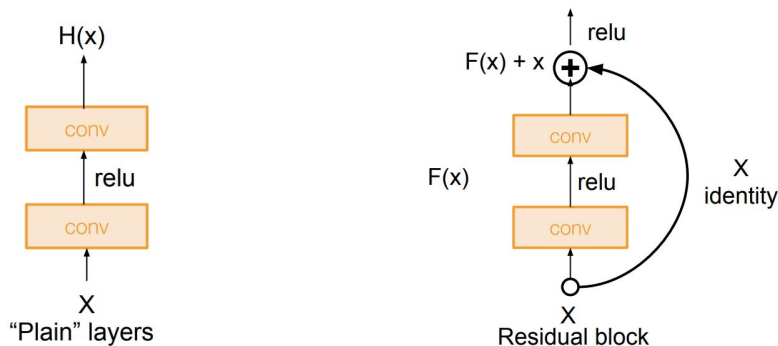
VGG-16 Conv3\_2



VGG-16 Conv5\_3

# Residual Networks (ResNets)

- Trend towards deeper networks
  - Makes sense to learn more low/mid/high-level features
  - But much harder to train
- ResNets allow for much deeper networks (~1000-layer networks)
- Have been found to work quite well on time series data



**Source:** He, Kaiming et al. "Deep Residual Learning for Image Recognition." 2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR) (2016): n. pag. Crossref. Web.

# External Resources

- Code base can be found at:
  - <https://github.com/valdivia4/Deep-Learning-Lunge-Detection>
- Implemented in Python
- Detailed documentation provided
  - No prior Python experience necessary